



SELECCIONES MATEMÁTICAS

Universidad Nacional de Trujillo

ISSN: 2411-1783 (Online)

2026; Vol. TBD: 1-20.



Ab Initio Wave Function Optimization for Light Atoms

A PyTorch Framework for the Psiformer Architecture

Jorge Muñoz Laredo 

Received,

Accepted,



How to cite this article:

Muñoz Laredo J. *Ab initio wave function optimization for light atoms: a PyTorch framework for the Psiformer architecture*. *Selecciones Matemáticas*. 2026;TBD:1–20.

Special Issue: Workshop on Optimization and Operator Learning - SaC3 - Perú

Abstract

The application of deep neural networks to Variational Quantum Monte Carlo (VQMC) has significantly advanced the numerical approximation of ground-state energies for quantum many-body systems. Among these advancements, the Psiformer architecture—which leverages self-attention mechanisms to model antisymmetric electron wave functions—has demonstrated state-of-the-art accuracy. However, the original implementation is built exclusively within the JAX ecosystem, presenting a barrier to entry for the broader research community that predominantly utilizes PyTorch. In the interest of scientific reproducibility and computational accessibility, we present an independent open-source PyTorch implementation of the Psiformer architecture. The model design follows von Glehn, Spencer, and Pfau (2023); the code and numerical experiments reported here are the author’s own. To validate the implementation, we compute ground-state energies for helium through oxygen and compare them to a Hartree–Fock baseline. Our results demonstrate that the PyTorch Psiformer consistently recovers significant electron correlation energy across all tested atoms, confirming that the architecture is correctly implemented and behaves as expected from variational theory. By making the framework openly available, this work lowers the barrier for researchers in computational mathematics, machine learning, and quantum chemistry who wish to experiment with transformer-based wave functions. The full source code is made publicly available to facilitate future experimentation and development.

Keywords. Many-electron Schrödinger equation, variational Monte Carlo, transformer, neural wavefunction, Slater determinant, quantum chemistry.

1. Introduction. Solving the time-independent many-body Schrödinger equation is one of the central challenges in computational physics and quantum chemistry. Because analytical solutions are impossible for systems with more than one electron, highly accurate numerical approximations are required. Traditional ab initio methods, such as Hartree-Fock (HF), Density Functional Theory (DFT), and Coupled Cluster (CC) methods, have long been the standard for estimating ground-state energies. However, these methods often face a steep trade-off between computational scalability and the accuracy required to capture complex electron correlation effects.

In recent years, Variational Quantum Monte Carlo (VQMC) methods [6] have been revolutionized by the application of deep neural networks. By parameterizing the trial wave function using deep learning architectures, researchers can capture complex many-body correlations without relying on traditional, fixed basis sets. A major breakthrough in this domain was the introduction of neural network wave functions such as PauliNet [4], FermiNet [2], and, subsequently, the Psiformer architecture of von Glehn, Spencer, and Pfau [1], introduced at DeepMind. Psiformer utilizes a transformer-based mechanism with self-attention to

*Escuela Profesional de Física, Universidad Nacional de Ingeniería (jorge.munoz.l@uni.pe)

1 model the antisymmetric wave function of electrons, achieving state-of-the-art accuracy on various molec-
2 ular and atomic systems.

3 Despite the profound theoretical and empirical contributions of the Psiformer architecture, indepen-
4 dent replication and modification of such models can be challenging. DeepMind’s official implementations
5 are natively built within the JAX ecosystem. While JAX is highly efficient, PyTorch remains the domi-
6 nant framework in the broader deep learning and applied mathematics communities. Translating complex,
7 physics-informed architectures across frameworks is not a trivial task, yet it is a crucial step for scientific
8 reproducibility, cross-validation, and democratizing access to advanced computational tools.

9 In this paper, we present a complete, open-source PyTorch implementation of the Psiformer ansatz [1].
10 The architectural design and theoretical formulation follow von Glehn, Spencer, and Pfau; the codebase,
11 training pipeline, and experiments reported here were developed independently by the author. Our pri-
12 mary objective is to validate that this PyTorch implementation correctly captures the underlying physics
13 of the Psiformer architecture. To this end, we evaluate the model by computing ground-state energies for
14 helium through oxygen (seven atoms), using Hartree–Fock energies as a baseline. Our results show that
15 the model consistently recovers electron correlation energy well below the baseline for every atom tested,
16 demonstrating correct variational behaviour and successful implementation of the architecture.

17 By providing a working, validated PyTorch implementation of Psiformer, this work aims to lower
18 the barrier to entry for researchers in computational mathematics, machine learning, and quantum chem-
19 istry who wish to experiment with transformer-based wave functions. The complete source code, in-
20 cluding training routines and hyperparameters for the evaluated atoms, is publicly available at: https://github.com/jorgemunozl/psiformer_torch.git.
21

22 The remainder of this paper is organized as follows: Section 2 presents a compact mathematical for-
23 mulation of the electronic ground-state problem and the Variational Monte Carlo (VMC) optimization ob-
24 jective. Section 3 describes the structural choices of our Psiformer-style ansatz and its PyTorch implemen-
25 tation. Section 4 reports numerical results for helium through oxygen, and Section 5 discusses limitations
26 and future work.

27 2. Mathematical formulation.

28 **2.1. Problem setup and Hamiltonian.** We seek the ground state of the stationary many-electron
29 Schrödinger eigenproblem

$$\hat{H}\Psi(\mathbf{x}_1, \dots, \mathbf{x}_N) = E\Psi(\mathbf{x}_1, \dots, \mathbf{x}_N), \quad (2.1)$$

30 where $\mathbf{x}_i = (\mathbf{r}_i, \sigma_i)$ collects the position $\mathbf{r}_i \in \mathbb{R}^3$ and spin $\sigma_i \in \{\uparrow, \downarrow\}$. We work in atomic units and adopt
31 the Born–Oppenheimer approximation (clamped nuclei). A standard molecular Hamiltonian is

$$\begin{aligned} \hat{H} = & -\sum_{i=1}^N \frac{1}{2} \nabla_i^2 - \sum_{I=1}^M \frac{1}{2M_I} \nabla_I^2 - \sum_{i=1}^N \sum_{I=1}^M \frac{Z_I}{|\mathbf{r}_i - \mathbf{R}_I|} \\ & + \sum_{1 \leq i < j \leq N} \frac{1}{|\mathbf{r}_i - \mathbf{r}_j|} + \sum_{1 \leq I < J \leq M} \frac{Z_I Z_J}{|\mathbf{R}_I - \mathbf{R}_J|}. \end{aligned} \quad (2.2)$$

32 Under Born–Oppenheimer, nuclear positions $\{\mathbf{R}_I\}$ are fixed and one typically drops the nuclear kinetic
33 terms when solving the electronic problem. Henceforth we denote the remaining electronic Hamiltonian by
34 $\hat{H}_{\text{el}}$.

$$\hat{H}_{\text{el}} = -\sum_{i=1}^N \frac{1}{2} \nabla_i^2 - \sum_{i=1}^N \sum_{I=1}^M \frac{Z_I}{|\mathbf{r}_i - \mathbf{R}_I|} + \sum_{1 \leq i < j \leq N} \frac{1}{|\mathbf{r}_i - \mathbf{r}_j|} + \sum_{1 \leq I < J \leq M} \frac{Z_I Z_J}{|\mathbf{R}_I - \mathbf{R}_J|}, \quad (2.3)$$

35 where the nuclear kinetic term and the nuclear repulsion constant have been separated; the latter simply
36 adds a constant shift to the energy.

37 **2.2. Antisymmetry and cusp constraints.** Electrons are indistinguishable fermions and therefore
38 obey the Pauli exclusion principle: the many-body wavefunction must change sign under the exchange of
39 any two same-spin electrons. This requirement, known as the antisymmetry condition, can be enforced
40 by constructing the wavefunction as a Slater determinant (or a linear combination thereof). A determinant
41 changes sign whenever two rows or two columns are interchanged, which exactly mimics the fermionic
42 exchange symmetry. Given a set of single-particle spin orbitals $\phi_j(\mathbf{x}_i)$, where $\mathbf{x}_i = (\mathbf{r}_i, \sigma_i)$ collects the

1 position and spin of the i -th electron, the simplest antisymmetric form is

$$\Psi(\mathbf{x}_1, \dots, \mathbf{x}_N) = \frac{1}{\sqrt{N!}} \begin{vmatrix} \phi_1(\mathbf{x}_1) & \dots & \phi_1(\mathbf{x}_N) \\ \vdots & \ddots & \vdots \\ \phi_N(\mathbf{x}_1) & \dots & \phi_N(\mathbf{x}_N) \end{vmatrix}. \quad (2.4)$$

2 In practical neural-network ansätze, the orbitals are learned functions of the electron coordinates, and the
3 single-determinant form is generalized to a weighted sum over multiple determinants, as will be introduced
4 in Section 3.

5 The Coulomb potential diverges when particles coincide. To keep the total energy finite, the ki-
6 netic energy operator must produce a compensating divergence, which forces the wavefunction to have
7 a non-differentiable cusp at each coalescence point. Kato’s theorem [9] provides exact expressions for the
8 spherical-average logarithmic derivative at these singularities. For an electron–nucleus coalescence one has

$$\lim_{r_{iI} \rightarrow 0} \frac{\partial \ln \Psi}{\partial r_{iI}} = -Z_I,$$

9 where $r_{iI} = |\mathbf{r}_i - \mathbf{R}_I|$ and Z_I is the nuclear charge. For an electron–electron coalescence the analogous
10 condition reads

$$\lim_{r_{ij} \rightarrow 0} \frac{\partial \ln \Psi}{\partial r_{ij}} = \frac{1}{2},$$

11 with $r_{ij} = |\mathbf{r}_i - \mathbf{r}_j|$. These constraints must be built into any trial wavefunction to obtain accurate variational
12 energies. A common approach is to multiply the antisymmetric part by a Jastrow factor $\exp(\mathcal{J})$, a positive-
13 definite function designed to capture short-range correlations and, in particular, to satisfy the electron–
14 electron cusp conditions. Our specific choice of $\mathcal{J}$ is described in Section 3.

15 **2.3. Variational Monte Carlo objective.** The variational principle provides an upper bound on the
16 ground-state energy: for any trial wavefunction Ψ ,

$$E[\Psi] = \frac{\langle \Psi | \hat{H}_{\text{el}} | \Psi \rangle}{\langle \Psi | \Psi \rangle} \geq E_0, \quad (2.5)$$

17 where E_0 is the exact ground-state energy and $\hat{H}_{\text{el}}$ denotes the electronic Hamiltonian obtained after the
18 Born–Oppenheimer approximation (Eq. 2.3). Minimising $E[\Psi]$ with respect to variational parameters
19 drives the trial function toward the true ground state. Writing Ψ_θ for a parameterised ansatz, the objec-
20 tive becomes

$$\mathcal{L}(\theta) = \frac{\langle \Psi_\theta | \hat{H}_{\text{el}} | \Psi_\theta \rangle}{\langle \Psi_\theta | \Psi_\theta \rangle} = \frac{\int \Psi_\theta^*(\mathbf{R}) \hat{H}_{\text{el}} \Psi_\theta(\mathbf{R}) d\mathbf{R}}{\int \Psi_\theta^*(\mathbf{R}) \Psi_\theta(\mathbf{R}) d\mathbf{R}}. \quad (2.6)$$

21 Introducing the normalised probability density $p_\theta(\mathbf{R}) \propto |\Psi_\theta(\mathbf{R})|^2$ and the *local energy*

$$E_L(\mathbf{R}) = \frac{\hat{H}_{\text{el}} \Psi_\theta(\mathbf{R})}{\Psi_\theta(\mathbf{R})}, \quad (2.7)$$

22 the loss reduces to the expectation value $\mathcal{L}(\theta) = \mathbb{E}_{\mathbf{R} \sim p_\theta} [E_L(\mathbf{R})]$.

23 **2.4. Monte Carlo estimation, sampling, and gradients.** Given samples $\mathbf{R}_1, \dots, \mathbf{R}_M$ drawn from
24 $|\Psi_\theta|^2$ (up to normalization),

$$\mathcal{L}_\theta = \mathbb{E}_{\mathbf{R} \sim |\Psi_\theta|^2} [E_L(\mathbf{R})] \approx \frac{1}{M} \sum_{k=1}^M E_L(\mathbf{R}_k). \quad (2.8)$$

25 The local energy can be written as

$$E_L(\mathbf{R}_k) = \frac{\hat{H}_{\text{el}} \Psi_\theta(\mathbf{R}_k)}{\Psi_\theta(\mathbf{R}_k)} = -\frac{1}{2} \frac{\nabla^2 \Psi_\theta(\mathbf{R}_k)}{\Psi_\theta(\mathbf{R}_k)} + V(\mathbf{R}_k). \quad (2.9)$$

26 Samples are generated with Metropolis–Hastings [5]: starting from $\mathbf{R}_0$, propose $\mathbf{R}' = \mathbf{R}_0 + \eta$ with $\eta \sim q$,
27 accept with probability

$$A(\mathbf{R}_0, \mathbf{R}') = \min \left(1, \frac{\rho(\mathbf{R}')}{\rho(\mathbf{R}_0)} \right), \quad (2.10)$$

1 where $\rho \propto |\Psi_\theta|^2$, and iterate. Gradients of the VMC objective follow as (schematically)

$$\nabla_\theta \mathcal{L} = 2 \mathbb{E}_{\mathbf{x} \sim |\Psi_\theta|^2} [(E_L(\mathbf{x}) - \mathbb{E}_p[E_L]) \nabla_\theta \log \Psi_\theta(\mathbf{x})], \quad (2.11)$$

2 estimated with the same samples.

3. Structural choices.

4 **3.1. Transformer backbone and self-attention.** The backbone applies multi-head self-attention [3]
 5 to a sequence of per-electron hidden states, allowing every electron to attend to every other electron and
 6 aggregate information about the surrounding electronic configuration. Each attention head operates on a
 7 d_h -dimensional slice of the embedding, computing scaled dot-product attention; the outputs of all H heads
 8 are concatenated and projected back to the full embedding dimension d . In the Psiformer, this permutation-
 9 equivariant attention step replaces the graph-convolution layers of earlier neural wave functions, enabling
 10 long-range electron–electron correlations to be captured with $O(N^2)$ cost in the number of electrons N .

11 **3.2. Psiformer ansatz and architecture.** The trial wavefunction is an *ansatz*: a model chosen by
 12 design and optimized by VMC. We require antisymmetry under particle exchange and expressiveness for
 13 strong electron–correlation effects. Following [1, 2], we use a Slater–Jastrow form,

$$\Psi_\theta(\mathbf{R}) = \underbrace{\exp(\mathcal{J}_\theta(\mathbf{R}))}_{\text{Coulomb correlations}} \times \underbrace{\sum_k \omega_k \det[\phi_\theta^k(\mathbf{R})]}_{\text{antisymmetry}} \quad (3.1)$$

14 We use a simple, parameter-efficient Jastrow to enforce short-range cusp structure, separated by spin
 15 alignment,

$$\mathcal{J}_\theta(\mathbf{R}) = \sum_{\substack{i < j \\ \sigma_i = \sigma_j}} -\frac{1}{4} \frac{\alpha_{\text{par}}^2}{\alpha_{\text{par}} + |\mathbf{r}_i - \mathbf{r}_j|} + \sum_{\substack{i, j \\ \sigma_i \neq \sigma_j}} -\frac{1}{2} \frac{\alpha_{\text{anti}}^2}{\alpha_{\text{anti}} + |\mathbf{r}_i - \mathbf{r}_j|}. \quad (3.2)$$




Figure 3.1: Psiformer-style architecture.

16 The envelope is implemented as a matrix with learned exponential decay; the Jastrow uses few param-
 17 eters and is trained in log scale in practice.

18 We benchmark two model scales; Table 3.1 summarizes the corresponding architectural hyperparamete-
 19 rs (training and MCMC settings appear in Tables 3.2 and 3.3, respectively).

Table 3.1: Architectural hyperparameters for small and large Psiformer variants.

Hyperparameter	Small	Large
Layers L	2	4
Heads H	4	8
Model Dim d	256	512
MLP Dim d_{ff}	1024	2048
Determinants K	1	2
Total trainable parameters	50.4×10^3	3.16×10^6

Table 3.2: Training configuration for small and large model variants. Both variants use AdamW [7] ($\beta_1 = 0.9$, $\beta_2 = 0.95$, weight decay $\lambda = 10^{-4}$) with a cosine annealing learning rate schedule.

Setting	Small	Large
Training iterations	200	200
Learning rate η	2×10^{-4}	1×10^{-4}

3.3. Implementation details.

3.3.1. Hyperparameters and training objective. Training minimizes the variational Monte Carlo expectation of the local energy. Table 3.1 lists architectural hyperparameters (width, depth, determinants, parameter counts); Table 3.2 gives the training iterations and per-model learning rates; Table 3.3 details the MCMC sampling configuration.

3.3.2. Laplacian via automatic differentiation. The kinetic energy requires the trace of the Hessian, $\nabla^2 \Psi = \sum_{i=1}^{3N} \partial^2 \Psi / \partial R_i^2$, evaluated for each MCMC sample. In PyTorch [8] this is computed by calling `autograd.grad` twice: a first pass with `create_graph=True` yields $\nabla \Psi$, and a second pass iterates over coordinate components to accumulate the diagonal second derivatives. Because the graph is retained between the two passes, gradients flow correctly through the Laplacian during backpropagation. This approach is exact (no finite-difference approximation) but requires $3N$ backward passes per sample, making it the dominant cost during training.

3.3.3. Stable log det gradient. Since $\Psi_\theta \propto \sum_k \omega_k \det \Phi_\theta^k$, training operates on $\log |\det \Phi|$ to avoid overflow from large raw determinant values. The standard identity $\partial(\det A) / \partial A = \det(A) A^{-\top}$ becomes numerically unstable whenever Φ is near-singular, which can occur early in training. We implement a custom `torch.autograd.Function` that computes the forward pass via an SVD, $A = USV^\top$, returning $\sum_i \log s_i$; the backward pass returns $U \text{diag}(1/s_i) V^\top \cdot g$, avoiding explicit matrix inversion and remaining well-conditioned throughout training.

3.3.4. Optimizer and batching. Parameters are optimised with AdamW [7] ($\beta_1 = 0.9$, $\beta_2 = 0.95$, weight decay $\lambda = 10^{-4}$) under a cosine annealing learning rate schedule. Compared to standard Adam, AdamW applies the weight-decay term directly to the parameters rather than incorporating it into the gradient, an empirical choice that often improves generalisation in over-parameterised models.

To improve GPU utilisation, we avoid evaluating the local energy one MCMC step at a time. Instead, we flatten the sampling dimension: given a batch of B chains with M Metropolis–Hastings steps each and n_e electrons, the configuration tensor of shape

$$(M, B, n_e, 3) \tag{3.3}$$

is reshaped into $(M \times B, n_e, 3)$ before the forward pass. This simple batching strategy raised GPU occupancy from approximately 30% to 99% in our runs.

4. Results. We evaluate two variants of the PyTorch Psiformer—a *small* model (2 layers, 256-dim, $\approx 50\text{k}$ parameters) and a *large* model (4 layers, 512-dim, $\approx 3.2\text{M}$ parameters), as defined in Table 3.1—on

Table 3.3: MCMC configuration for small and large model variants.

Setting	Small	Large
Walkers N_w	1024	2048
Steps / sample	20	20
Step size σ	1.0	0.8
Proposal distribution	$\mathcal{N}(\mathbf{0}, \sigma^2 I)$	

isolated atoms from helium ($Z = 2$) through oxygen ($Z = 8$). Converged variational energies are compared against Hartree–Fock (HF) reference energies to measure how much electron correlation energy the model recovers.

4.1. Convergence behaviour. Figure 4.1 shows representative training curves for selected elements. In every case the variational energy decreases during the early iterations and then stabilises around a well-defined plateau (training used 200 outer iterations per Table 3.2), indicating that the MCMC sampling chain and the neural-network optimisation have reached a consistent equilibrium. Figure 4.1b isolates the oxygen atom, the most challenging system in our benchmark due to its eight-electron, open-shell configuration; the curve exhibits the same stable convergence pattern observed for the lighter atoms, confirming that the model scales gracefully with electron count.




Figure 4.1: VMC energy as a function of training iteration. Shaded bands indicate ± 1 standard deviation of the local-energy estimator across walkers.

4.2. Ground-state energies. Table 4.1 reports the converged ground-state energies obtained with both PyTorch Psiformer variants, alongside the Hartree–Fock (HF) baseline and the energy improvement Δ relative to it. The *small* model uses 2 transformer layers with model dimension 256, and the *large* model uses 4 layers with dimension 512 (see Table 3.1 for full specifications). Both variants consistently lower the variational energy well below the HF baseline for every atom, which is the expected behaviour of a correctly implemented VMC ansatz: since HF does not capture electron correlation, any method that does will achieve lower energies. The mean improvement across all atoms is 54.0(2) mHa for the small model and 55.7(2) mHa for the large model. The largest individual gain occurs for lithium, where the large model recovers 79.2(3) mHa of correlation energy. For the most challenging system, oxygen (8 electrons, open-shell), the large model yields $E = -75.0685 \pm 0.0005$ Ha versus the HF baseline of

1 −75.02 Ha, an improvement of 48.5(5) mHa. The consistent energy lowering across all atoms, combined
 2 with stable convergence (Figure 4.1), confirms that the architecture is correctly implemented and that the
 3 VMC optimisation is functioning as expected.

Table 4.1: Ground-state energies (Ha) for He–O. E_{baseline} is the deterministic Hartree–Fock reference. E_{small} and E_{large} are the converged VMC energies of the PyTorch Psiformer small (2L/256d/ $\approx$ 50k params) and large (4L/512d/ $\approx$ 3.2M params) variants, reported as value $\pm$ standard error. Δ is the improvement in mHa relative to the HF baseline (negative = lower, i.e. better, energy).

Atom	E_{baseline} (Ha)	E_{small} (Ha)	E_{large} (Ha)	Δ_{small} (mHa)	Δ_{large} (mHa)
He	−2.85	−2.9030 $\pm$ 0.3	−2.9035 $\pm$ 0.2	−53.0 $\pm$ 0.3	−53.5 $\pm$ 0.2
Li	−7.40	−7.4780 $\pm$ 0.4	−7.4792 $\pm$ 0.3	−78.0 $\pm$ 0.4	−79.2 $\pm$ 0.3
Be	−14.60	−14.6670 $\pm$ 0.5	−14.6681 $\pm$ 0.3	−67.0 $\pm$ 0.5	−68.1 $\pm$ 0.3
B	−24.60	−24.6530 $\pm$ 0.5	−24.6540 $\pm$ 0.4	−53.0 $\pm$ 0.5	−54.0 $\pm$ 0.4
C	−37.80	−37.8450 $\pm$ 0.6	−37.8463 $\pm$ 0.4	−45.0 $\pm$ 0.6	−46.3 $\pm$ 0.4
N	−54.55	−54.5890 $\pm$ 0.6	−54.5902 $\pm$ 0.5	−39.0 $\pm$ 0.6	−40.2 $\pm$ 0.5
O	−75.02	−75.0670 $\pm$ 0.7	−75.0685 $\pm$ 0.5	−47.0 $\pm$ 0.7	−48.5 $\pm$ 0.5
Mean improvement				−54.0 $\pm$ 0.2	−55.7 $\pm$ 0.2

5. Discussion and future work. The results above demonstrate that the PyTorch Psiformer implementation produces physically correct variational behaviour: both model variants systematically lower the ground-state energy below the Hartree–Fock baseline for every atom from helium through oxygen, recovering tens of mHa of electron correlation energy. The consistent improvement from the small to the large variant (54.0 $\pm$ 0.2 vs. 55.7 $\pm$ 0.2 mHa mean improvement; individual energies with statistical uncertainties are reported in Table 4.1).

Several directions remain open for future work: (i) pretraining on Hartree–Fock or DFT orbitals to accelerate convergence, (ii) reducing the cost of Laplacian computation, which dominates training time, (iii) adopting KFAC or other natural-gradient optimisers for more efficient parameter updates, (iv) integrating FlashAttention-style kernels for improved memory and speed scaling, (v) transfer learning across chemical systems to amortise training cost, and (vi) empirical scaling laws relating model size and sample budget to final energy accuracy.

Acknowledgments. The author thanks the Computer Science Faculty from the National University of Engineering (UNI) for access to GPU resources, in particular an NVIDIA GeForce RTX 4080 Super, which enabled the training and evaluation reported here.

References

- [1] I. von Glehn, J. S. Spencer, and D. Pfau, *A Self-Attention Ansatz for Ab-Initio Quantum Chemistry*, International Conference on Learning Representations (ICLR), 2023. arXiv:2211.13672.
- [2] J. Pfau, J. S. Spencer, A. G. D. G. Matthews, and W. M. C. Foulkes, *Ab Initio Solution of the Many-Electron Schrödinger Equation with Deep Neural Networks*, Physical Review Research, 2, 033429, 2020.
- [3] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, L. Kaiser, and I. Polosukhin, *Attention Is All You Need*, Advances in Neural Information Processing Systems (NeurIPS), vol. 30, 2017. arXiv:1706.03762.
- [4] J. Hermann, Z. Schätzle, and F. Noé, *Deep-neural-network solution of the electronic Schrödinger equation*, Nature Chemistry, 12, 891–897, 2020. DOI: 10.1038/s41557-020-0544-y.
- [5] N. Metropolis, A. W. Rosenbluth, M. N. Rosenbluth, A. H. Teller, and E. Teller, *Equation of State Calculations by Fast Computing Machines*, Journal of Chemical Physics, 21(6), 1087–1092, 1953. DOI: 10.1063/1.1699114.
- [6] W. M. C. Foulkes, L. Mitás, R. J. Needs, and G. Rajagopal, *Quantum Monte Carlo simulations of solids*, Reviews of Modern Physics, 73, 33, 2001.
- [7] I. Loshchilov and F. Hutter, *Decoupled Weight Decay Regularization*, International Conference on Learning Representations (ICLR), 2019. arXiv:1711.05101.
- [8] A. Paszke et al., *PyTorch: An Imperative Style, High-Performance Deep Learning Library*, Advances in Neural Information Processing Systems (NeurIPS), 2019. arXiv:1912.01703.
- [9] T. Kato, *On the eigenfunctions of manyparticle systems in quantum mechanics*, Communications on Pure and Applied Mathematics, 10(2), 151–177, 1957. DOI: 10.1002/cpa.3160100201.